

Computer Vision

Group 12 | 0786031 廖俊凱 0516044 陳思妤 0856733 黃明翰

Introduction

Image stitching, in simple terms, is for an input there should be a group of images, the output is a composite image such that it is a culmination of image scenes. At the same time, the logical flow between the images must be preserved.

For example, think about sea horizon while you are taking few photos of it. From a group of these images, we are essentially creating a single stitched image, that explains the full scene in detail. It is quite an interesting algorithm.


Let's first understand the concept of image stitching. Basically, if you want to capture a big scene and your camera can only provide an image of a specific resolution and that resolution is 640 by 480, it is certainly not enough to capture the big panoramic view. So, what we can do is to capture multiple images of the entire scene and then put all bits and pieces together into one big image. Such photos of ordered scenes or collections are called panoramas. The entire process of acquiring multiple image and converting them into such panoramas is called as image stitching. And finally, we have one beautiful big and large photograph of the scenic view.

The basic geometry of the problem is well understood, and consists of estimating a 3×3 camera matrix or homography for each image. This estimation process needs an initialization, which is typically provided by user input to approximately align the images, or a fixed image ordering. In this assignment, we formulate stitching as a multi-image matching problem, and use invariant local features to find matches between all of the images.



Implementation Procedure

- A. Interest points detection & feature description by SIFT
1. Take 16x16 square window around detected feature
 2. Compute edge orientation for each pixel
 3. Throw out weak edges (threshold gradient magnitude)
 4. Create histogram of surviving edge orientations
 5. 16 cells * 8 orientations = 128 dimensional descriptors



```
1. def sift(input_img, img_name):
2.     img = np.uint8(input_img)
3.     descriptor = cv2.xfeatures2d.SIFT_create()
4.     (kps, des) = descriptor.detectAndCompute(img, None)
5.     cv2.drawKeypoints(img, kps, img, (0,255,255), flags=cv2.DRAW_MATCHES_FLAGS
        _DRAW_RICH_KEYPOINTS)
6.     cv2.imwrite('./results/sift/sift_{}.jpg'.format(img_name.split('.')[0]),im
        g)
7.     return kps, des
```

B. Feature matching by SIFT features

1. Exact and match SIFT features between all of the images
2. SIFT features are located at scale-space maxima/minima of a difference of Gaussian function
3. At each feature location, a characteristic scale and orientation is established
4. Give a similarity-invariant frame in which to make measurements

Once features have been extracted from all n images (linear time), they must be matched. Since multiple images may overlap a single ray, each feature is matched to its k nearest neighbors in feature space (we use $k = 4$). This can be done in $O(n \log n)$ time by using a k -d tree to find approximate nearest neighbors. A k -d tree is an axis aligned binary space partition, which recursively partitions the feature space at the mean in the dimension with highest variance.

```
1. def find_match(des1, des2):
2.     dist_matrix = np.zeros((len(des1), len(des2)))
3.     # find the distance between each and every descriptors
4.     for index1 in range(len(des1)):
5.         for index2 in range(len(des2)):
6.             dist = np.linalg.norm(des1[index1] - des2[index2])
7.             dist_matrix[index1][index2] = dist
8.             #sort the vector to get the two best match with the smallest distance
9.             #match = [descriptor's number(descriptor i), the two best match with descriptor i, distance with the first match, distance with the second match]
10.            match = []
11.            for i in range(dist_matrix.shape[0]):
12.                two_best_match = np.argsort(dist_matrix[i])[:2]
13.                match.append([i, two_best_match, dist_matrix[i][two_best_match[0]], dist_matrix[i][two_best_match[1]]])
14.            #perform ratio distance with ratio = 0.75
15.            final_match = []
16.            for i in range(len(match)):
17.                if match[i][2] < match[i][3] * 0.75:
18.                    final_match.append([match[i][0], match[i][1][0]])
19.            return final_match
20.
21. def draw_match(img1, img2, kps1, kps2, match, img_name):
22.     height1 = img1.shape[0]
23.     width1 = img1.shape[1]
24.     height2 = img2.shape[0]
25.     width2 = img2.shape[1]
26.     output_img = np.zeros((max(height1, height2), width1+width2, 3), dtype = 'uint8')
27.     output_img[0:height1, 0:width1] = img1
28.     output_img[0:height2, width1:] = img2
29.
30.     for index1, index2 in match:
31.         color = list(map(int, np.random.randint(0, high=255, size=(3,))))
32.         pts1 = (int(kps1[index1].pt[0]), int(kps1[index1].pt[1]))
33.         pts2 = (int(kps2[index2].pt[0] + width1), int(kps2[index2].pt[1]))
34.         cv2.line(output_img, pts1, pts2, color, 1)
35.     cv2.imwrite('./results/match/matchline_'+img_name+'.jpg', output_img)
```

C. RANSAC to find homography matrix H

1. Randomly select a RANSAC sample from the sample set, that is, 4 matching point pairs
2. Calculate the homography matrix H according to these 4 matching point pairs
3. According to the sample set, the homography matrix H, and the error metric function, calculate the consistent set consensus that satisfies the current the homography matrix H and return the number of elements in the set
4. Determine whether the optimal (maximum) consistent set is based on the number of elements in the current consistent set, and if so, update the current optimal consistent set
5. Update the current error probability p, if p is greater than the minimum allowed error probability, repeat (1) to (4) to continue iterating until the current error probability p is less than the minimum error probability

```
1. def homomat(min_match_count: int, src, dst):
2.     A = np.zeros((min_match_count * 2, 9))
3.     # construct the two sets of points
4.     for i in range(min_match_count):
5.         src1, src2 = src[i, 0, 0], src[i, 0, 1]
6.         dst1, dst2 = dst[i, 0, 0], dst[i, 0, 1]
7.         A[i * 2, :] = [src1, src2, 1, 0, 0, 0, -src1 * dst1, -src2 * dst1, -dst1]
8.         A[i * 2 + 1, :] = [0, 0, 0, src1, src2, 1, -src1 * dst2, -src2 * dst2, -
9.             dst2]
10.
11.     # compute the homography between the two sets of points
12.     [_, S, V] = np.linalg.svd(A)
13.     m = V[np.argmax(S)]
14.     m *= 1 / m[-1]
15.     H = m.reshape((3, 3))
16.     return H
17.
18. def ransac(final_match, kps_list, min_match_count, num_test: int, threshold: float):
19.     if len(final_match) > min_match_count:
20.         src_pts = np.array([kps_list[1][m[1]].pt for m in final_match]).reshape(-
21.             1, 1, 2)
22.         dst_pts = np.array([kps_list[0][m[0]].pt for m in final_match]).reshape(-
23.             1, 1, 2)
24.         min_outliers_count = math.inf
25.         while(num_test != 0):
26.             indexs = np.random.choice(len(final_match), min_match_count, replace=False)
27.             homography = homomat(min_match_count, src_pts[indexs], dst_pts[indexs])
28.             # Warp all left points with computed homography matrix and compare SSDs
29.             src_pts_reshape = src_pts.reshape(-1, 2)
30.             one = np.ones((len(src_pts_reshape), 1))
31.             src_pts_reshape = np.concatenate((src_pts_reshape, one), axis=1)
32.             warped_left = np.array(np.matmul(homography) * np.matmul(src_pts_reshape).T)
33.             for i, value in enumerate(warped_left.T):
34.                 warped_left[:, i] = (value * (1 / value[2])).T
35.             # Calculate SSD
36.             dst_pts_reshape = dst_pts.reshape(-1, 2)
```


```

37.         dst_pts_reshape = np.concatenate((dst_pts_reshape, one), axis=1)
38.         inlier_count = 0
39.         inlier_list = []
40.         for i, pair in enumerate(final_match):
41.             ssd = np.linalg.norm(np.array(warped_left[:, i]).ravel() - dst_pts_resh
ape[i])
42.             if ssd <= threshold:
43.                 inlier_count += 1
44.                 inlier_list.append(pair)
45.
46.         if (len(final_match) - inlier_count) < min_outliers_count:
47.             min_outliers_count = (len(final_match) - inlier_count)
48.             best_homomat = homography
49.             best_matches = inlier_list
50.             num_test -= 1
51.         return best_homomat, best_matches
52.     else:
53.         raise Exception("Not much matching keypoints exists!")

```

D. Warp image to create panoramic image

Once we have established a homography, i.e. we know how the second image will look from the current image's perspective, we need to transform it into a new space. That is, the slightly distorted, and altered image that we see from our periphery. This process is called warping. We are converting an image, based on a new transformation.



```

1. def forward_warp(h, w, warped_img, img_grid):
2.     res = np.zeros((h, w, 3), dtype='uint8')
3.     for x, y, im in zip(warped_img[0], warped_img[1], img_grid):
4.         res[math.floor(y)+top, math.floor(x)+left] = img2[im[1], im[0], :]
5.     cv2.imwrite('./results/warp/warp_res/forward_'+img_name+'.jpg', res)
6.     blend(res, 'forward_')
7.
8. def inverse_warp(h, w, corners, homography):
9.
10.    res_nn = np.zeros((h, w, 3), dtype='uint8')
11.    res_bi = np.zeros((h, w, 3), dtype='uint8')
12.
13.    # Create image 2 grid in image 1 coordinate
14.    b, t, r, l = math.ceil(max(corners[:, 1])), math.floor(min(corners[:, 1])), math.ceil(
max(corners[:, 0])), math.floor(min(corners[:, 0]))
15.
16.    img2_trans_grid = [[n, m, 1] for n in range(l, r) for m in range(min(t, 1), b)]
17.
18.    # Inverse mapping points on image 1 to image 2
19.    img2_trans_inv = np.array(np.mat(np.linalg.inv(homography)) * np.mat(img2_trans_gri
d).T)

```

```

20.     img2_trans_inv /= img2_trans_inv[2]
21.
22.     for x, y, im in zip(img2_trans_inv[0], img2_trans_inv[1], img2_trans_grid):
23.         if math.ceil(y) < img2.shape[0] and math.ceil(y)>0 and math.ceil(x) < img2.shap
e[1] and math.ceil(x)>0:
24.             # Nearest Neighbor
25.             res_nn[im[1]+top, im[0]+left] = img2[int(y + 0.5), int(x + 0.5), :]
26.
27.             # Bilinear interpolation
28.             res_bi[im[1]+top,im[0]+left] = (img2[math.ceil(y), math.ceil(x), :]*((y-
math.floor(y))*(x-math.floor(x)))+
29.                                             img2[math.floor(y), math.floor(x), :]*
((math.ceil(y)-y)*(math.ceil(x)-x))+
30.                                             img2[math.ceil(y), math.floor(x), :]*(
(y-math.floor(y))*(math.ceil(x)-x))+
31.                                             img2[math.floor(y), math.ceil(x), :]*(
(math.ceil(y)-y)*(x-math.floor(x))))
32.             cv2.imwrite('./results/warp/warp_res/backward_nn_'+img_name+'.jpg', res_nn)
33.             cv2.imwrite('./results/warp/warp_res/backward_bi_'+img_name+'.jpg', res_bi)
34.             #blend(res_nn, 'inverse_nn_', 0.5)
35.             #blend(res_bi, 'inverse_bi_', 0.5)
36.             return linear_blend(res_bi, window_size)
37.
38. def warp(homography):
39.
40.     # Transform image 2 with homography matrix
41.     img2_grid = [[n, m, 1] for n in range(img2.shape[1]) for m in range(img2.shape[0])]
42.
43.     img2_trans = np.array(np.mat(homography) * np.mat(img2_grid).T)
44.     img2_trans /= img2_trans[2]
45.
46.     # Find transformed four corners of image 2 on image 1 coordinate system
47.     corners = np.zeros((4, 3))
48.     for p, im in zip(img2_trans.T, img2_grid):
49.         if im[0] == 0 and im[1] == 0:
50.             corners[0] = p
51.         elif im[0] == 0 and im[1] == img2.shape[0] - 1:
52.             corners[1] = p
53.         elif im[0] == img2.shape[1] - 1 and im[1] == 0:
54.             corners[2] = p
55.         elif im[0] == img2.shape[1] - 1 and im[1] == img2.shape[0] - 1:
56.             corners[3] = p
57.
58.     # Blended image size
59.     global top,left,bottom,r
60.     top = max(0, math.ceil(-min(corners.T[1])))
61.     bottom = max(img1.shape[0], math.ceil(max(corners.T[1])))+top
62.     left = max(0, math.ceil(-min(corners.T[0])))
63.     right = max(img1.shape[1], math.ceil(max(corners.T[0])))+left
64.     r = max(img1.shape[1], math.ceil(min(corners[2][0],corners[3][0])))+left
65.
66.     #Forward warping
67.     #forward_warp(bottom, right, img2_trans, img2_grid)
68.
69.     # Inverse warping
70.     return inverse_warp(bottom, right, corners, homography)

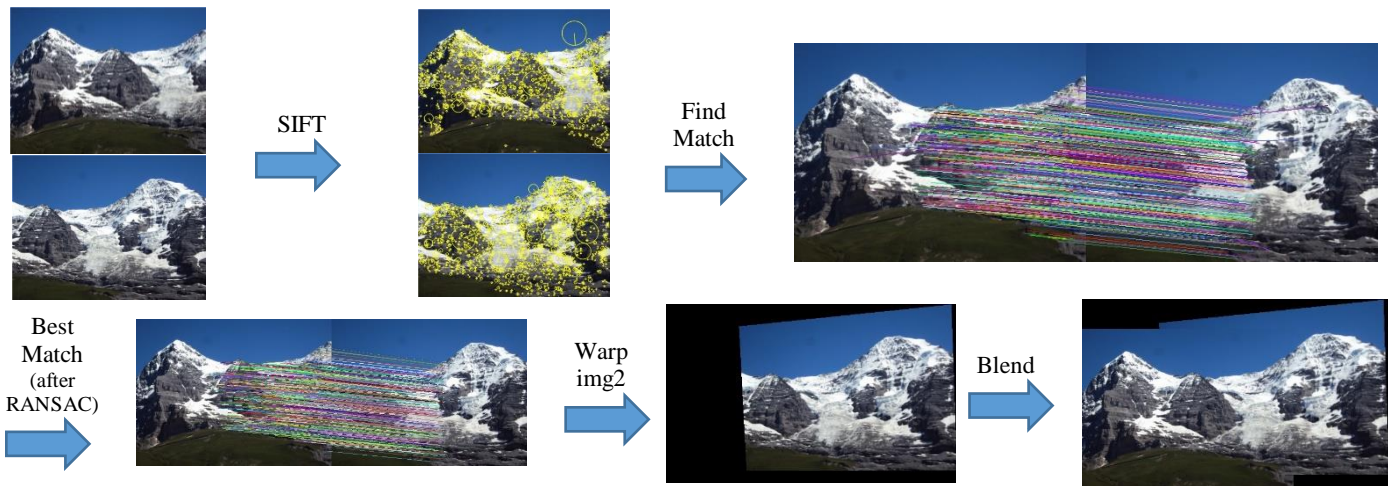
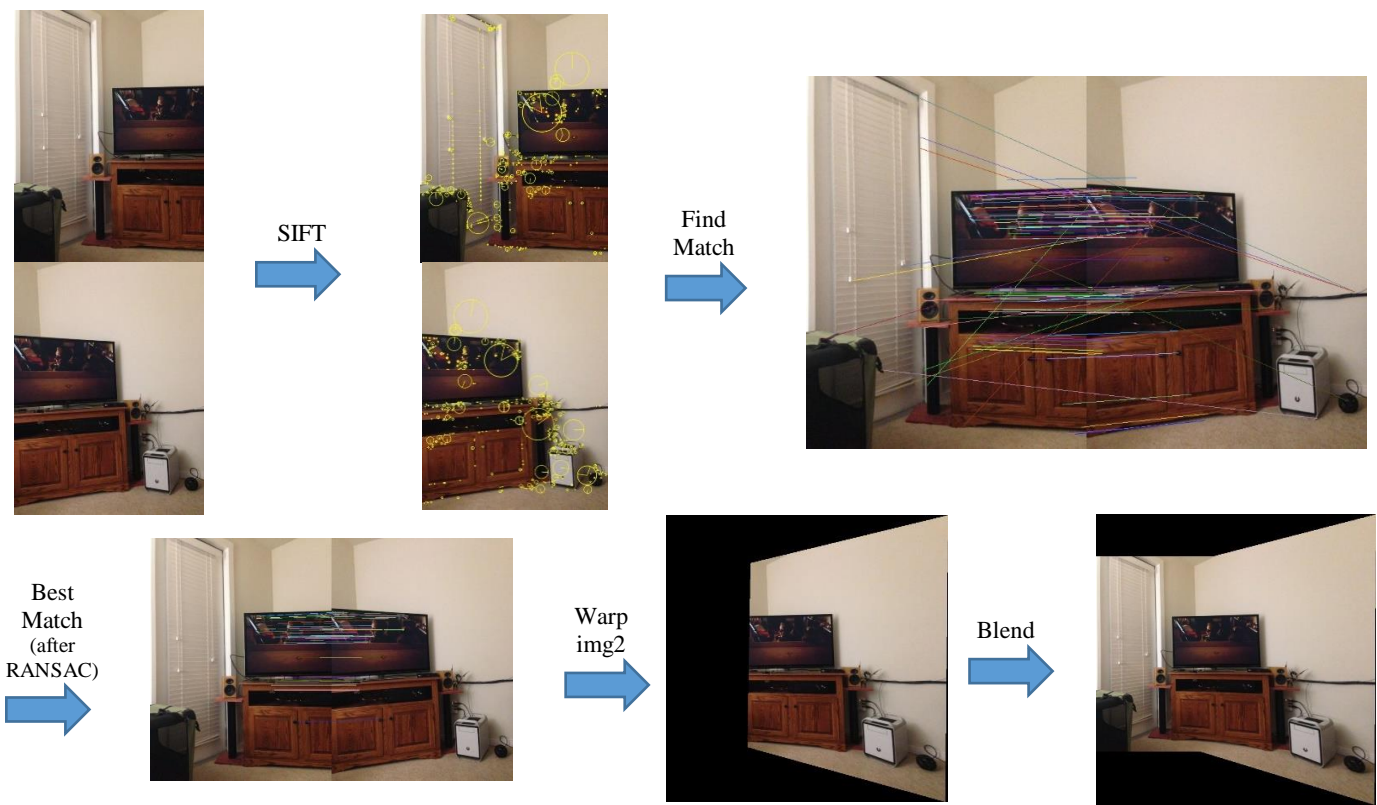
```

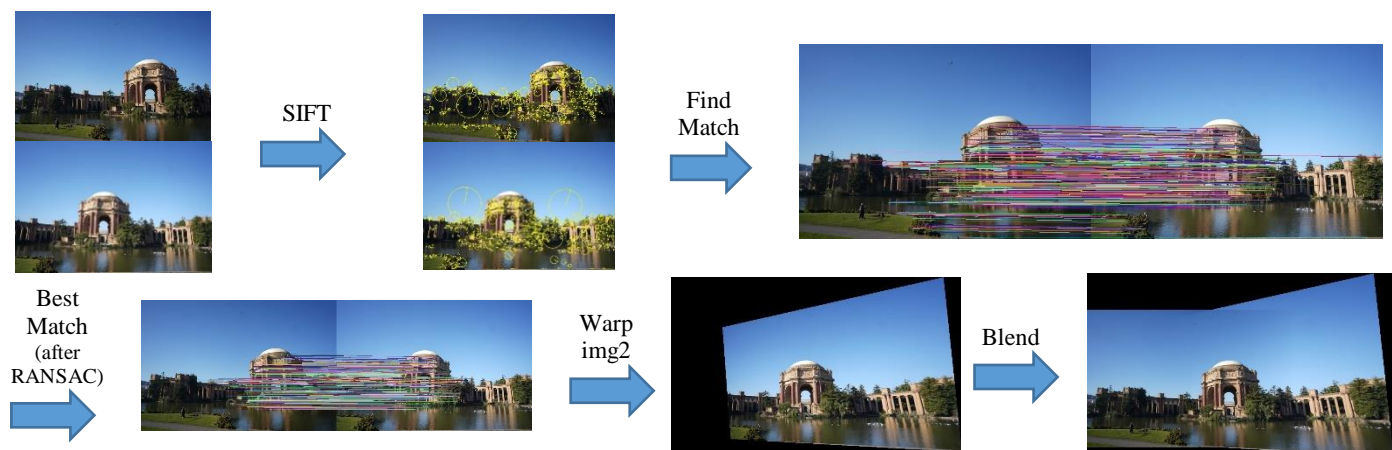
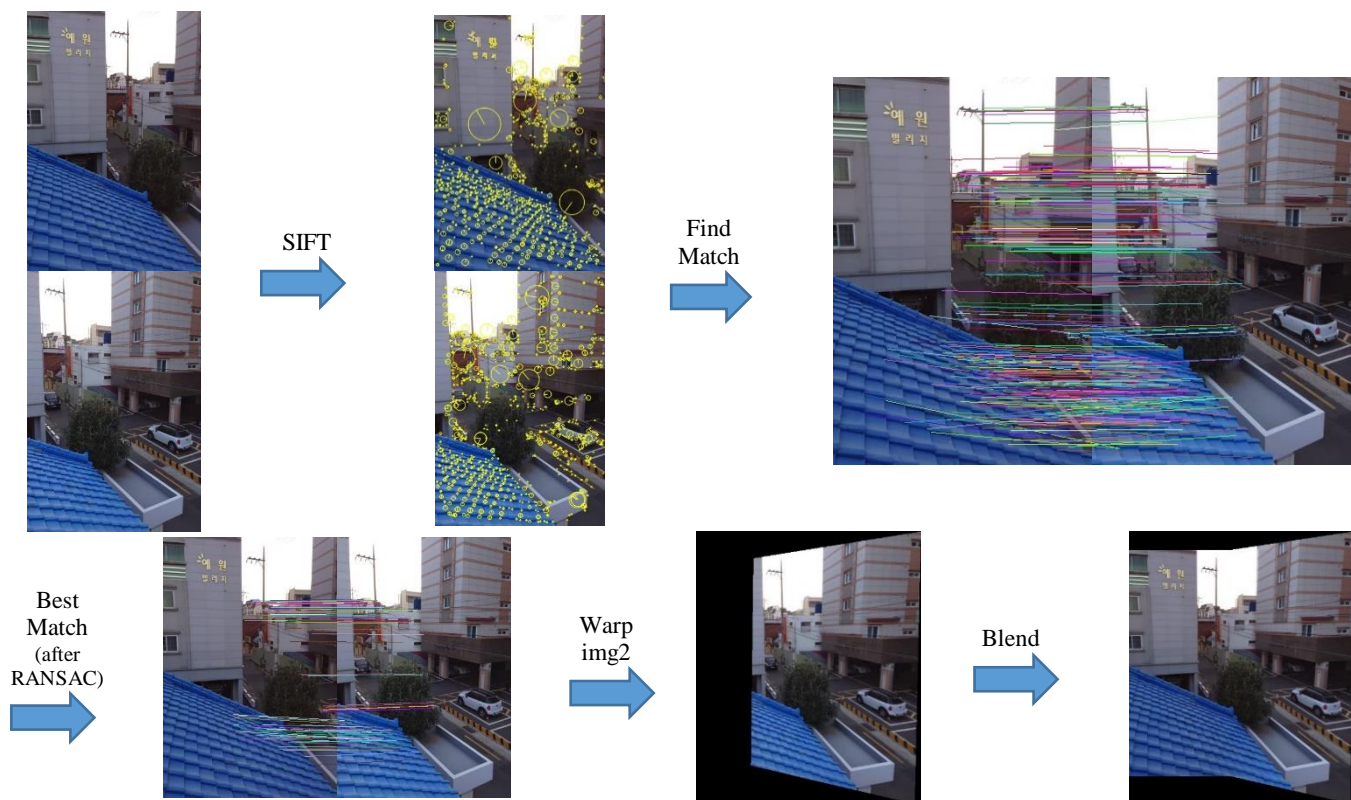
E. Blending function

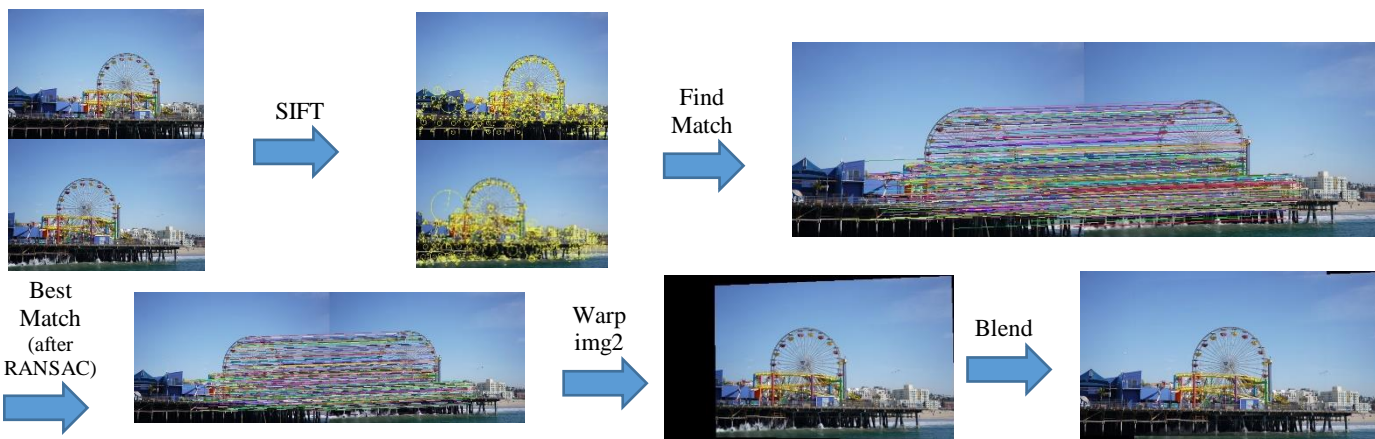
Ideally each sample (pixel) along a ray would have the same intensity in every image that it intersects, but in reality, this is not the case. Even after gain compensation some image edges are still visible due to a number of unmodelled effects, such as vignetting (intensity decreases towards the edge of the image), parallax effects due to unwanted motion of the optical center, mis-registration errors due to mismodelling of the camera, radial distortion and so on. Because of this, a good blending strategy is important. Here we use the linear blending technique with a flexible window size.

```
1. def linear_blend(res, window_size=0):
2.     res = res.copy()
3.     alpha = step_a = window = step_w = 0
4.     for m in range(img1.shape[1]):
5.         alpha += step_a
6.         window += step_w
7.         for n in range(img1.shape[0]):
8.             if sum(res[n+top, m:left]) != 0:
9.                 if window==0:
10.                     step_w = 1/(img1.shape[1]-m)
11.                     window+=step_w
12.                 if window > window_size and alpha==0:
13.                     step_a = 1/(img1.shape[1]-m)
14.                     alpha+=step_a
15.                 res[n+top, m:left] = alpha * res[n+top, m:left] + (1 - alpha) * img1[n,
16. m]
17.             else:
18.                 res[n+top, m:left] = img1[n, m]
19.         cv2.imwrite('./results/warp/linear_window_warp_'+img_name+'.jpg', res)
20.     return res
21.
22. def blend(res, name, alpha=0.5):
23.     res = res.copy()
24.     for m in range(img1.shape[1]):
25.         for n in range(img1.shape[0]):
26.             if sum(res[n+top, m:left]) != 0:
27.                 res[n+top, m:left] = alpha * res[n+top, m:left] + (1 - alpha) * img1[n,
28. m]
29.             else:
30.                 res[n+top, m:left] = img1[n, m]
31.         cv2.imwrite('./results/warp/warp_'+name+img_name+'.jpg', res)
```




Experiment Result







Cropped Result (Big Picture)






Discussion


1. Forward & Inverse Warping

A. Forward Warping

Directly transform image 2 to the coordinate system of image 1. If the result is not an integer, take the nearest integer as result. Since the warping points are often non-integer sample, many samples “o” are **not assigned values**.



B. Inverse Warping



Using inverse warping, we can avoid the problem in forward warping since every point in the new coordinate system (image1) has a value. The different interpolation technology will get different result. We use **bilinear interpolation** for better result.



Nearest-Neighbor Interpolation (TOP) looks more **cruffy** than Bilinear Interpolation (BOT)

2. Blending

There are many ways to blend image1 and image2:


A. Naïve Blending

For $\alpha \in [0,1]$, new image = $\alpha * \text{image1} + (1 - \alpha) * \text{image2}$. The result is not good since there are obvious edges in the intersection of two images.




B. Linear Blending

Linearly feather the intersection part of two image. The result is better now since the edges disappear. But there still some problem.



C. Linear Blending with a window

Take a closed look at the pole in the result of the linear blending, there is a **ghosting problem** since two images are not perfectly matched. We can narrow the blending part by using a proper window to solve this problem.



3. Comparison with CV2 built-in function

In this assignment, we were manually performing keypoint detection, feature matching, giving us access to the homography matrix used to warp our two input images into a panorama.

And while OpenCV's built-in `cv2.createStitcher` and `cv2.Stitcher_create` functions are certainly capable of constructing accurate, aesthetically pleasing panoramas, one of the primary drawbacks of the method is that it abstracts away any access to the homography matrices.

One of the assumptions of real-time panorama construction is that the scene itself is not changing much in terms of content. Once we compute the initial homography estimation we should only have to occasionally recompute the matrix.

Not having to perform a full-blown key-point matching and RANSAC estimation gives us a tremendous boost of speed when building our panorama, so without access to the raw homography matrices, it would be challenging to take OpenCV's built-in image stitching algorithm and convert it to real-time.

Conclusion

In this assignment, we learned how to perform multiple image stitching using OpenCV and Python. Using both OpenCV and Python we were able to stitch multiple images together and create panoramic images. First, we implement feature matching with interest points detection and feature description by SIFT. Then, we warp images to create panoramic image after establishing homography matrix H according to RANSAC. Finally, our output panoramic images were not only accurate in their stitching placement but also aesthetically pleasing as well.

Reference

1. <https://www.pyimagesearch.com/2018/12/17/image-stitching-with-opencv-and-python/>
2. <https://kushalvyas.github.io/stitching.html>
3. <http://matthewalunbrown.com/papers/ijcv2007.pdf>
4. <https://blog.csdn.net/zhangjunhit/article/details/83088627>

Work Assignment Plan Between Team Members

0516044 陳思妤	SIFT + Feature Matching
0786031 廖俊凱	RANSAC + Homography
0856733 黃明翰	Wrap + Blending